RESEARCH AND DIAGNOSIS · 2 SEPTEMBER 2026 · JOAT FOR EVEN REALITIES G2

The assistant is possible. The current brain is not the reason it feels impossible.

Nine days of the brain's own log, a live benchmark of its reply model on this Mac, a read of the voice path in both packages, and Even's primary sources all point the same way: the 30-second replies, the three-word answers, and the "I don't know" habit come from four specific choices, not from the goal.

THE VERDICT

Nothing about the G2 blocks a capable, conversational assistant. Even ships a supported "bring your own agent" switch for its built-in assistant, the Hub SDK gives a plugin the mic, the display, and open network access, and at least a dozen community projects already run Claude or other frontier models on this exact hardware.

What is broken is the brain's own pipeline. A 3-billion-parameter local model is told to refuse anything not in its fact block, so it refuses jokes and Hamlet. A text-level wake-word gate discards roughly 85 percent of what you say. A one-second silence detector chops a request into fragments, and only the fragment containing "JOAT" is answered. Every character of reply is painted as image tiles over Bluetooth at half a second to two seconds per tile.

The recommendation: stand up one small "JOAT agent" service backed by a frontier Claude model and your existing tools, point Even AI's Agent Configure at it this week, and then fix the four faults in the Hub plugin so JOAT's own surface catches up using the same brain.

MEASURED · BRAIN.LOG, 25 AUG TO 2 SEP, 93,027 LINES

What the log says happened to your voice

The brain logs no transcript text by design, so the numbers below are counts of outcomes. Every spoken utterance ends in exactly one of these rows.

Outcome of an utterance	Count	What it means